
EvoMan - Framework 1.0

DOCUMENTATION

KARINE MIRAS

KARINE.SMIRAS@GMAIL.COM

August 27, 2019

Contents

1	Introduction	2
2	EvoMan framework	2
2.1	First steps	2
2.2	Platform structure	2
2.2.1	Controlling agents	2
2.2.2	Training objectives	6
2.2.3	Sensors	6
2.2.4	Running an experiment	7
2.3	Classes	8
2.3.1	Environment class	8
2.3.2	Controller class	13
2.4	Recovery	13
3	Dissertation summary	14
3.1	Objectives	14
3.2	Methodology	15
3.3	Results	15
4	Demos	18
4.1	Dummy demo	18
4.2	Human demo	18
4.3	Specialist agent	19
4.4	Generalist agent	19
4.5	Individual evolution	19
4.6	Multi evolution	19
4.7	Coevolution	19

1 Introduction

EvoMan framework is a platform for developing/testing optimization algorithms, and was developed in the Federal University of ABC in Santo Andre / Sao Paulo / Brazil during the research of a masters project.

This document provides details of how to install and use the platform. Furthermore, a very brief overview of baseline experiments is also presented. The results of more baseline experiments can be found in [Miras, 2016b] and [Miras, 2016a].

2 EvoMan framework

2.1 First steps

Follow the steps below to install and start using the EvoMan framework.

- `python -m pip install -U pygame --user`
- `pip install tmx`
- `pip install numpy`
- `git clone https://github.com/karinemiras/evoman_framework`
- run one demo: `python controller_specialist_demo.py`

You should be able to see an EA+ANN playing the game against multiple opponents. Some examples can be seen here <https://youtu.be/ZqaMjd1E4ZI>.

2.2 Platform structure

EvoMan provides **8** predator-prey games (Fig. 1) inspired on the electronic game **MegaMan II**.

2.2.1 Controlling agents

Both prey (player) and predator (enemy) may be controlled by an Artificial Intelligence **algorithm** (agent), or assume its **default** behaviour.

For the **player**, the default behaviour is waiting for commands from some input device. For the **enemy**, the default behaviour is an static attack composed by fixed rules, and this is different for each enemy. Details about the dynamics of each game (how the enemy attacks, and how the environment affects the player) can be found in [Miras, 2016b].




Figure 1: Games

Combining the controlling methods for the agents (enemy and player) provides us interesting simulation **modes**, as shown in Figure 2.




Figure 2: Simulation modes

- (a) **No-evolution:** the player is controlled by a human through an input device, while the enemies present static attacks defined as fixed rules.
- (b) **Individual evolution:** the player is controlled by a user-defined algorithm, and the enemies by static attacks defined as fixed rules.

- (c) **Co-evolution**: both player and enemy are controlled by user-defined algorithms.
- (d) **Co-learning**: the player is controlled by a human through an input-device, and the enemy by a user-defined algorithm.

2.2.2 Training objectives

When training an agent, it is possible to choose between **single** or **multiple** objectives (Fig. 3).

If you choose single objective, you will be training an **specialist** agent against one enemy. If you choose multiple objectives, you will be training a **generalist** agent, fighting against more than one enemy in the same training phase.



Figure 3: Training objectives

2.2.3 Sensors

The agents have access to **20 sensors** (Fig. 4). They are: 16 for distances from the player to the projectiles; 2 for distances from the player to the enemy; 2 for concerning the direction the agents are facing.




Figure 4: Sensors for agents

2.2.4 Running an experiment

Using the framework for running experiments can be done as depicted in Figure 5. The optimization algorithm must search for solutions to the problem (single or multiple objectives), and every time a solution is tested, the game loop is run.

During the game loop, agent controllers receive the sensors state at each game time step, and take decisions for all their possible actions.




Figure 5: Experiment stream

2.3 Classes

2.3.1 Environment class

Environment is the main class of the framework. It may be instantiated as it is or implemented as needed.

When instantiating the class with no new implementations and providing no parameters, it will use default fitness functions, and and default configurations.

```
from environment import Environment
env = Environment()
env.play()
```

Figure 6: Environment class simple instantiation

Parameters

The list below describes all parameters for this class. All parameters assume their default values when not provided during the instantiation.

- **multiplemode**: whether objective is multiple or single. The multiple objectives mode required informing multiple games in the 'enemies' parameter. It may assume [yes] for multiple mode, or [no] for single mode. Default value is [no].
- **enemies**: games list. Must be an array from 1 to 8 positions containing numbers from 1 to 8. Default value is [8].
- **loadplayer**: excludes player agent from the game, useful for testing. May assume [yes] for keeping player, or [no] for excluding player. Default value is [yes].
- **loadenemy**: excludes enemy agent from the game, useful for testing. May assume [yes] for keeping player, or [no] for excluding player. Default value is [yes].
- **level**: integer number for defining the game level of difficulty. The greater the level, the harder the game is, i.e., the more projectiles one agent must hit the other with to defeat it. Default value [1].
- **playermode**: mode for controlling player. May be [human] for controlling by keyboard/joystick, or [ai] for controlling by an algorithm. Default value [ai]. If you do not have a joystick, use the arrows to move, space to jump, and shift to shoot.
- **enemymode**: mode for controlling enemy. May be [static], controlled by fixed rules, or [ai], controlled by an algorithm. Default value [ai].
- **speed**: game run velocity. May be [normal] for 30 frames a second (real time) or [fastest] for the maximum speed possible in pygame. Default value [normal].¹

¹The option 'fastest' did not work for me on Mac OS Sierra. In case it also does not work for

- **inputscoded**: applies several transformations to agent sensors, increasing the number of variables. May be [yes] for coding, or [no] for not coding. Default value [no].
- **randomini**: positions the enemy at random initial positions on the platform in each game run. May be [yes] for varying position, or [no] for fixed position. Default value [no].
- **sound**: game sound. May be [on] for sound on, or [off] for sound off. Sound only works in player human mode. Default value [on].
- **contacthurt**: defines which agent is injured during body contact. May be [player] for injuring the player, or [enemy] for injuring the enemy. Default value [player].
- **logs**: May be [on] for printing logs, or [off] for not printing logs. Default value [on].
- **savelogs**: turns on the exportation of logs to the file evoman_logs.txt. May be [yes] for exporting, or [no] for not exporting. Default value [yes].
- **timeexpire**: integer number for defining the game run maximum duration. Useful for avoiding unnecessary long game runs. The greater the time, the longer the run is. Default value [3000].
- **overturetime**: integer number for defining the interval time for before/after starting/finishing a game. The overture only happens in player human mode. The greater the time, the longer the overture takes. Default value [100].
- **clockprec**: precision of time measurement. Pygame's time counter is not totally accurate, which means games will not present exactly the same environmental conditions every run. This happens because the time counting in pygame is used for changing the game frames, and that depends of the conditions of the operational system. This is useful for simulating a dynamic environment, which helps to achieve more robust solutions. This value may

you, and you wish to run experiments using the fastest speed of pygame, you may use a Linux virtual machine.

be [low] for using the low precision function `clock.tick()` function, or [medium] for using the medium precision function `clock.tick_busy_loop()`. Default value [low].

- **player_controller**: a Controller class object representing the controller for the player agent. Default value is an object of the original Controller class, where a default method is implemented taking random actions for the agent.
- **enemy_controller**: a Controller class object representing the controller for the enemy agent. Default value is an object of the original Controller class, where a default method is implemented taking random actions for the agent.

Methods

- **update_solutions**(self, solutions): receives a list containing the experiment results found so far and saves it in the current game state. The user may store anything in it.
- **update_parameter**(self, name, value): any parameter must be updated using this method. 'name' is the name of the parameter, and 'value' is the new value it will assume.
- **get_num_sensors**(self): returns the number of sensors configured for the controllers.
- **save_state**(self, name): saves the current game state, containing solutions and configured parameters in the files `evoman_solstate_`'name' and `evoman_paramstate_`'name'.txt respectively, in which 'name' is the name to be given to the state.
- **load_state**(self, name): loads a saved game state as current, setting parameters and solutions backup, where 'name' is the name of the state saved.
- **get_playerlife**(self): returns the number of energy points kept by the player in the match.
- **get_enemylife**(self): returns the number of energy points kept by the enemy in the match.

- **get_time(self)**: returns the duration of the match (in game steps).
- **play(self, pcont=None, econt=None)**: executes a game run using the solutions given as 'pcont'(player) and 'econt'(enemy) for the controllers. When using the default controllers these parameters are not necessary, as agents will take random actions.

When in single objective mode, it returns: fitness, player life, enemy life, and game run time, in this order as single values. When in multi objective mode, it returns: fitness, player life, enemy life, and game run time, in this order as 4 lists containing values of each game (lists are ordered by the order of enemies parameter provided).

- **fitness_single(self)**: this method should be **implemented** and replaced to define a different fitness function for training agents in single objective mode. The default fitness function is expressed by Equation 1.

$$fitness = \gamma * (100 - e_e) + \alpha * e_p - \log t \quad (1)$$

where e_e , e_p are the energy measures of enemy and player respectively, varying from 0 to 100; t is number of time steps of game run; and constants γ and α assume values 0.9 and 0.1 respectively.

- **cons_multi(self, values)**: this method should be **implemented** and replaced to define a different fitness function for training agents in multiple mode. The default function is expressed by Equation 2.

This method must consolidate all fitnesses of single objective solutions into one only fitness. 'values' is a list with the fitnesses calculated by `fitness_single(self)` method for every game configured in the experiment.

$$fitness' = \bar{f} - \sigma \quad (2)$$

where $fitness'$ is the consolidated $fitness$ by \bar{f} and σ , that are the mean and the standard deviation of fitnesses against the m enemies chosen for training.

2.3.2 Controller class

This is the class to be implemented when defining a controller structure.

The `control(self, params, cont)` method should be to **implemented** to define the agent controller. If you do not implement this method, the agent will use the default method that takes random actions.

The 'params' param received is a numpy array containing all the 20 sensors' values, to be read at each game time step.

When implementing control methods for the controllers, the returned variables at each game time step must be Boolean, and the number of variables may vary for each agent:

- The player controller may take **actions**: walk left, walk right, jump, shoot, and release of the jump.²
- The enemies 1, 2, 3 and 4 controllers may take **4** actions, defined as attack 1 until attack 4.
- The enemy 5 and 6 controller may take **3** actions, defined as attack 1 until attack 3.
- The enemy 7 and 8 controller may take **6** actions, defined as attack 1 until attack 6.

2.4 Recovery

If you need to stop the simulation for any reason, your experiments can be recovered from the latest successful generation. The demos present examples of how to recover. Furthermore, use the script `loop.sh` to run your experiment in infinite loop. This is important because from time to time pygame simply “dies” with an error that (apparently) can not be caught.

²The control should return the actions in this order.

3 Dissertation summary

Title: Neuroevolution for the construction of a generic strategy with EvoMan environment.

University library link for dissertations and theses (in portuguese):
<http://biblioteca.ufabc.edu.br/index.html>

3.1 Objectives

Problem: generalizing learning for multiple problems

Application: electronic games

General Video Game Playing (GVGP) is a recent research field that aims creating autonomous generic controllers able to play different games.




Figure 7: GVGP

Objective 1 - Creating a testing environment that permits training autonomous agent controllers using electronic games in different simulation modes.

Objective 2 - Testing NeuroEvolution algorithms to create specialist (single game) and generalist (generalist) agents using EvoMan.

3.2 Methodology

For assessing solutions found by the optimization algorithm, it is necessary to define a success measure, usually called fitness.

The fitness defined for single objective problem was the default one. Also for multi objective problem, a default fitness was used. The greater the fitness, the better the solution is, i.e., maximization.

The optimization method used was NeuroEvolution, that trains and defines neural networks using evolutionary algorithms. The method was applied in two ways:

- Evolving only the network weights for a perceptron and a multilayer (10 and 50 neurons) network: using a Genetic Algorithms and the LinkedOpt algorithm.
- Evolving the network weights and topology using the NEAT algorithm.

The **success** criterion for the solutions was:

- **Effectiveness** means to zero the energy (life) of the enemy.
- **Efficiency** is the amount of energy preserved by the player.

Solutions can be **classified** as:

- **Bad solutions** do not win the game (unsuccessfull)
- **Medium solutions** win the game preserving less than 50 points of energy (successfull)
- **Good solutions** win the game preserving more than 50 points of energy (successfull)

3.3 Results

For specialist solutions, some enemies presented greater challenge than others, but all enemies were shown to be beatable (Fig. 8).




Figure 8: Algorithms performance by game for single objective solutions.

For generalist solutions, several groups of enemies were composed and tested. From the total of 8 enemies, all (28) possible pairs of enemies were combined and tested using the NEAT algorithm as a reference, and from the best pairs, some triplet were also tested.

To realize a final generalization test for generic agent controllers, a measure called Gain was used, as shown in Figure 9. The measure goes from -800 to 800, and the greater it is, the better is the generalization power of the agent.

$$g = \sum_{i=1}^n p_i - \sum_{i=1}^n e_i$$

Figure 9: Gain measure

where e is the energy of the enemy, from 0 to 100; p is the energy of the player, from 0 to 100; $n=8$, is the total number of games in the test.

The best groups obtained by NEAT were:

- pairs: (2, 4) (2, 6) (7, 8)

- trios: (2, 6, 5) (6, 5, 1) (2, 5, 1) (6, 7, 4)

These best groups were tested using the other 2 algorithms also. The best final training group found was: 7 and 8. (Fig 11)




Figure 10: Algorithms performance by game for multiple objective solutions.

The group reached the greatest Gain value (40), by the Genetic Algorithm with a neural network having a hidden layer with 10 neurons. (Fig. 10)

Also this group defeated the greatest number of enemies (5), by the Genetic Algorithm with a perceptron neural network.

	NEAT			GAP			AGM10			AGM50			LOPT			LOPTM10			LOPTM50		
	E	T	V	E	T	V	E	T	V	E	T	V	E	T	V	E	T	V	E	T	V
1	1			1	1		1	1		1	1		1	1		1	1		1	1	
2			67	2		15	2		74	2		8	2		24	2		77	2		
3				3			3			3			3			3			3		
4				4			4			43			4			4			4		25
5			6	5		45	5		60	5		48	5		52	5		41	5		61
6				6		2	6			6			6			6			6		
7		42		7	53		7			7	19		7	70		7	68		7		
8	60			8	56		8	30		8	45		8	18		8	50		8	26	

Figure 11: Algorithms performance by enemy for multiple objective solutions, final energy measures.

Best generalist solutions videos:

youtu.be/w97qitczuL8

youtu.be/vQ_xCShU1gQ

youtu.be/ic68VCjMIMI

4 Demos

4.1 Dummy demo

File: `dummy_demo.py`

Runs one game configured for ai player using random actions as controller, and a static enemy.

4.2 Human demo

File: `human_demo.py`

Runs one game configured for human player and static enemy.

4.3 Specialist agent

File: controller_specialist_demo.py

Runs solutions found by an EA for each one of the enemies separately. Presents examples of good and bad players.

4.4 Generalist agent

File: controller_generalist_demo.py

Runs best solution found by an EA, trained using as multi objective.

4.5 Individual evolution

File: optimization_specialist_demo.py

Starts an experiment evolving an agent against one single static enemy, using an EA.

4.6 Multi evolution

File: optimization_generalist_demo.py

Starts an experiment evolving an agent against multiple static enemies, using an EA.

4.7 Coevolution

optimization_coevolution_demo.py

Starts an experiment co-evolving player and enemy using EAs for both.

References

[Miras, 2016a] Miras, Karine, F. O. (2016a). An electronic-game framework for evaluating coevolutionary algorithms. *arXiv preprint arXiv:1604.00644*.

[Miras, 2016b] Miras, Karine, F. O. (2016b). Evolving a generalized strategy for an action-platformer video game framework. In *2016 IEEE Congress on Evolutionary Computation (CEC)*, pages 1303–1310. IEEE.